


No-pressure result across two MoE architectures


M1 Max · n=5 admitted runs per arm · model-specific unseen prompts · exact outputs

Upstream 16
Upstream 96
ElasticUMA 96/hot16

(a) Decode throughput



(b) Peak physical footprint



Bars show medians. Qwen uses 256 generated tokens; Gemma uses 128.