


Committed ElasticUMA versus untouched upstream fixed caches


M1 Max · Qwen3.6-35B-A3B Q4 · n=5 admitted runs per arm · identical outputs within each protocol

Upstream 16
Upstream 96
ElasticUMA 96/hot16

(a) Decode throughput



(b) Peak physical footprint



Bars show medians; dots show all measured runs. Source: admitted V6/V7 receipts.